



KubeCon

CloudNativeCon

Europe 2025



**C.A.L.L.I.N.G. now
I'm calling you,
calling you now**

terra tauri
Staff Software Engineer - Grafana Labs

Mario Macías
Staff Software Engineer - Grafana Labs



KubeCon



CloudNativeCon

Europe 2025



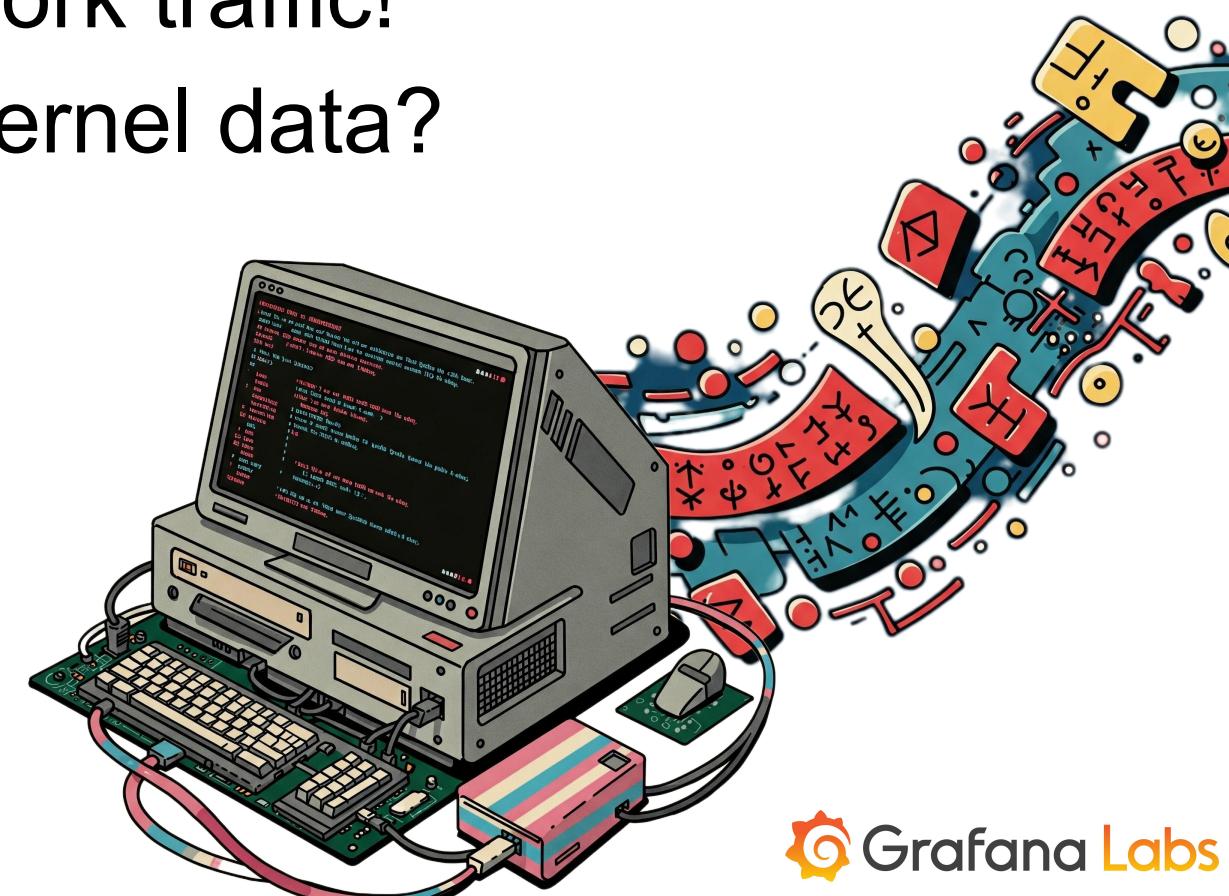
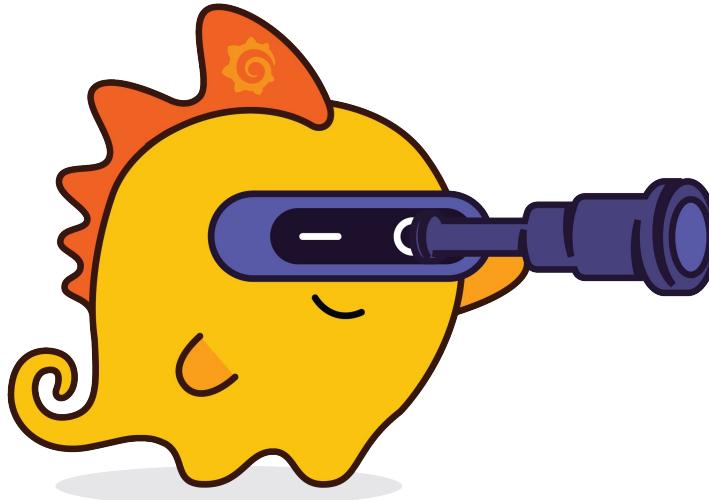
Today You'll Learn About

- eBPF
- Breaking Kubernetes API
- How Beyla enriches kernel data
- Scaling eBPF enrichment



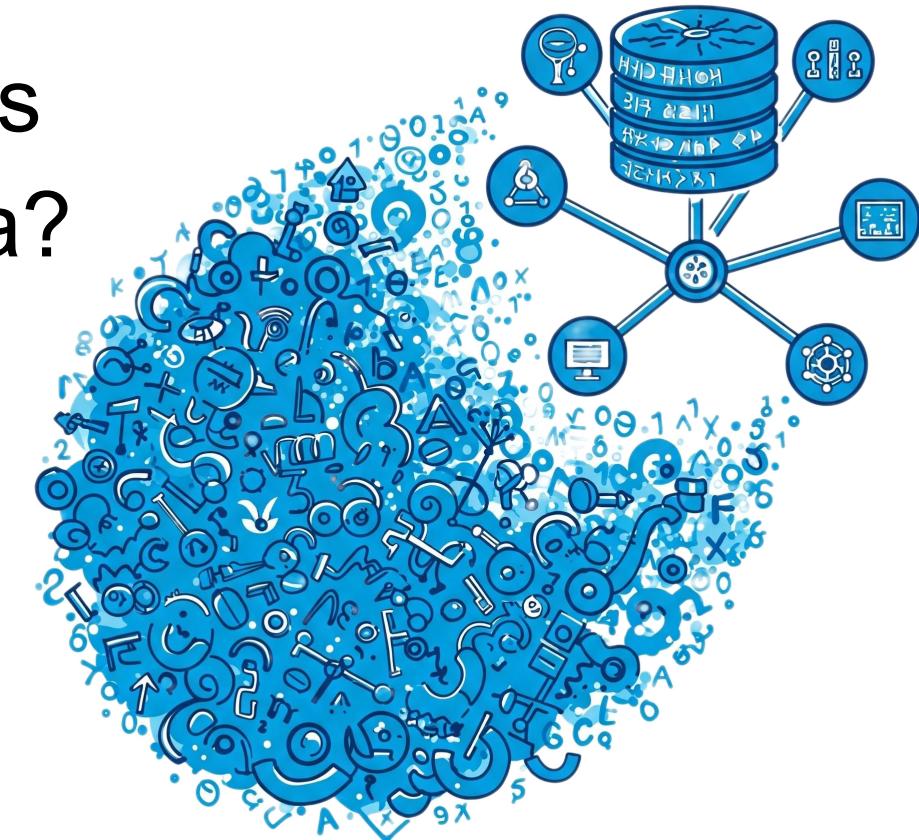
eBPF: overview

- Enables safe and efficient extension of the kernel
- For performance, We use hooks into `tc`
- Allows us to snoop on network traffic!
- - ... but how useful is raw kernel data?



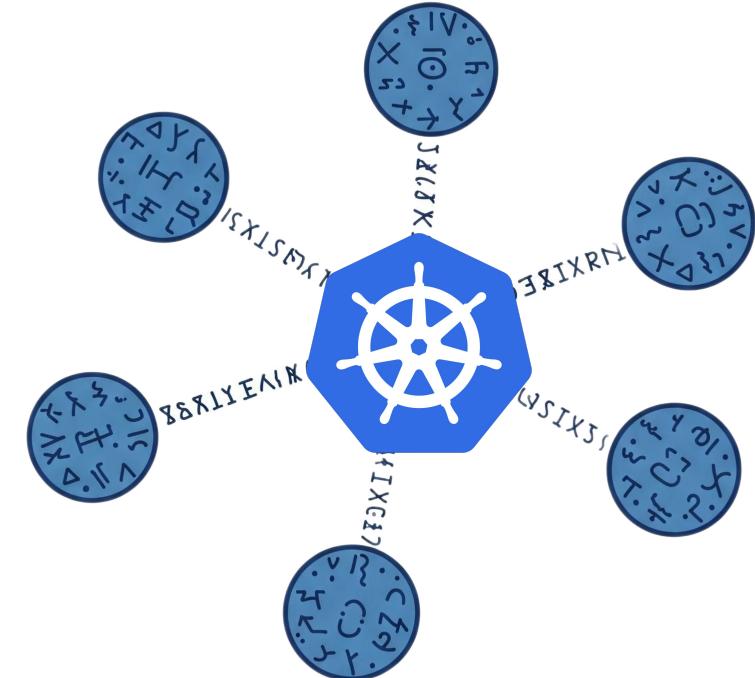
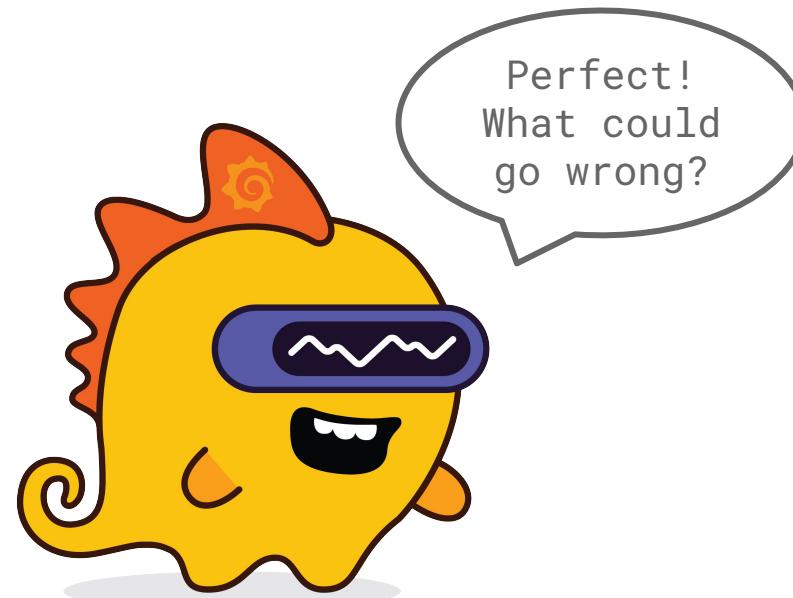
Kernel Data is Raw

- Each Node has a kernel
- Kernel data gives IP addresses
- In Kubernetes, we care about Pods
- How can we enrich our kernel data?



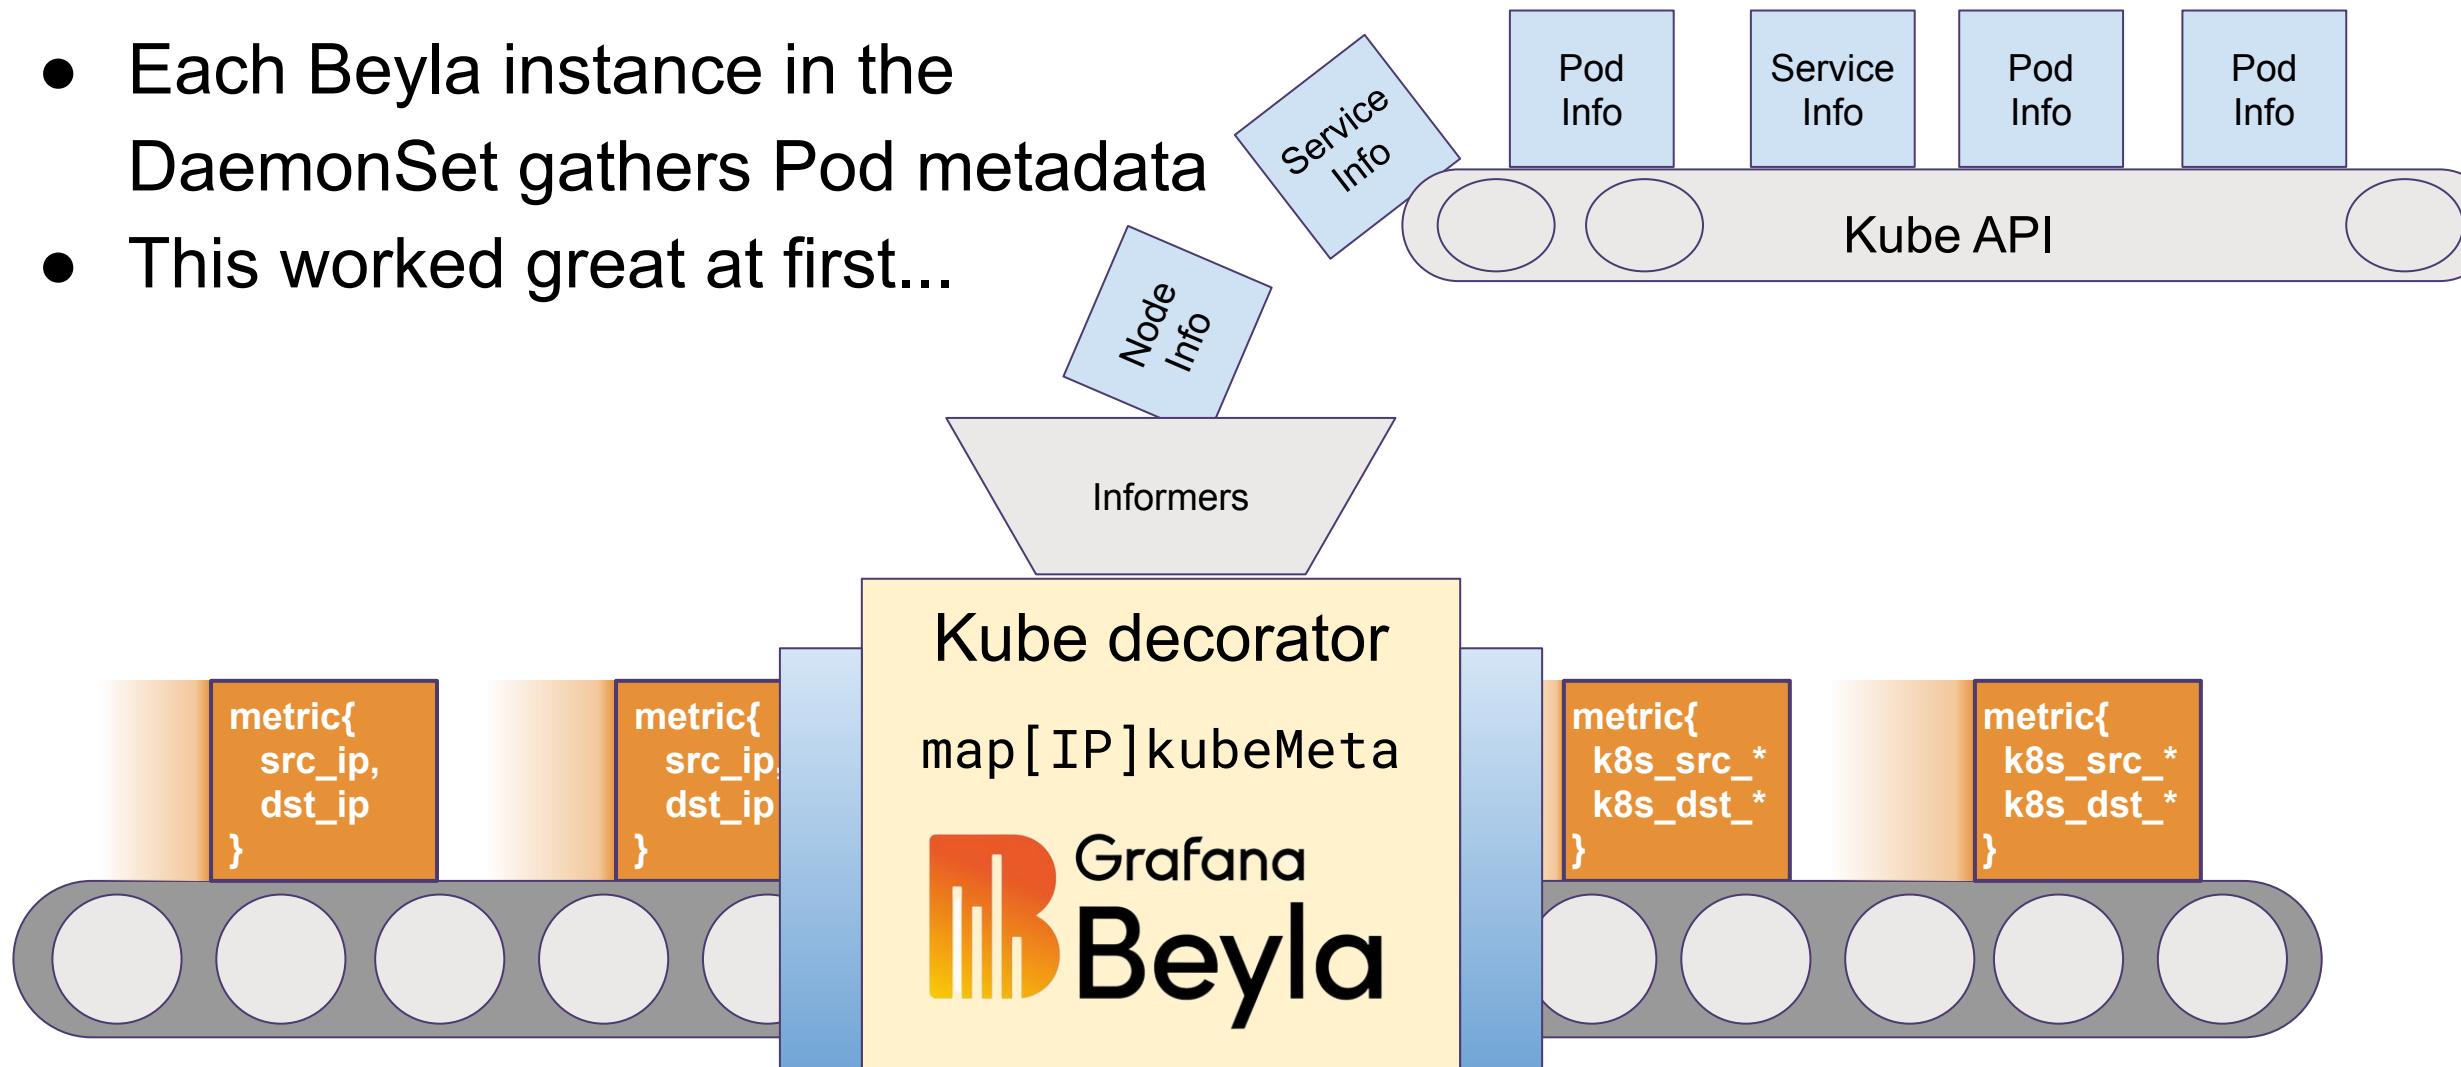
eBPF on Kubernetes

- DaemonSet = 1 pod per kernel
- We have this very tempting API
- It knows everything about the cluster



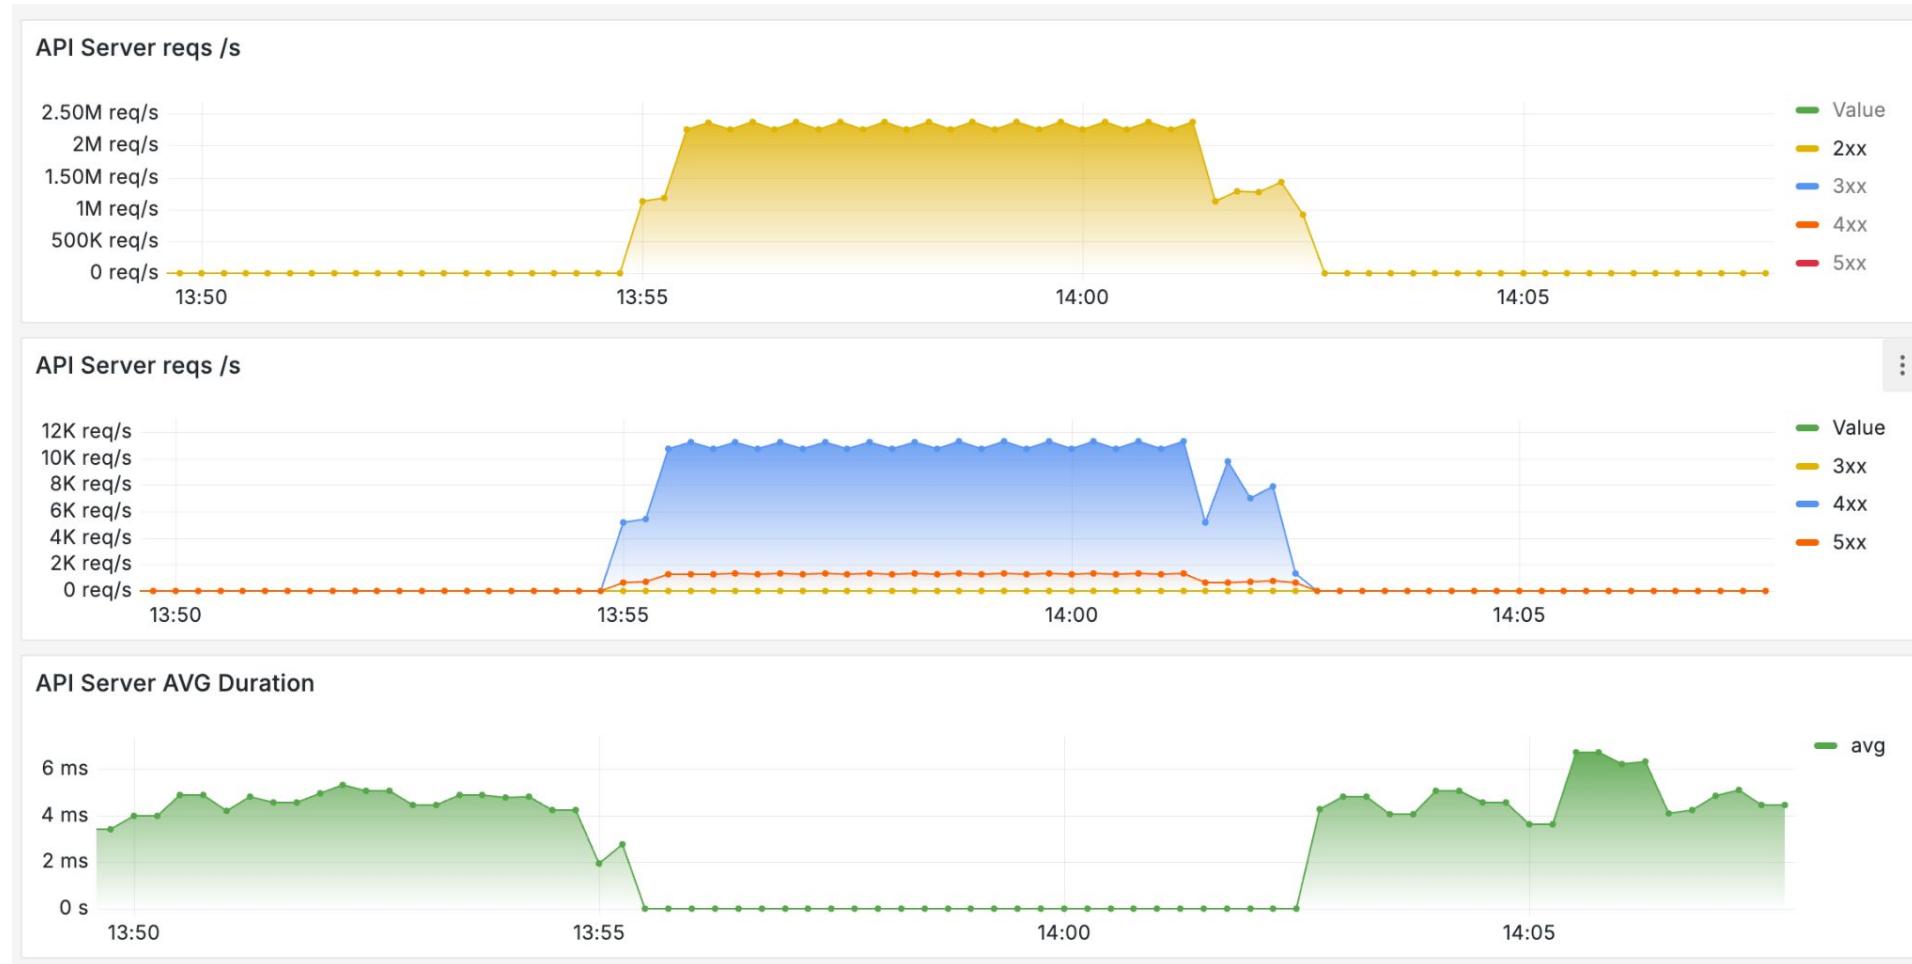
Our first attempt

- Each Beyla instance in the DaemonSet gathers Pod metadata
- This worked great at first...

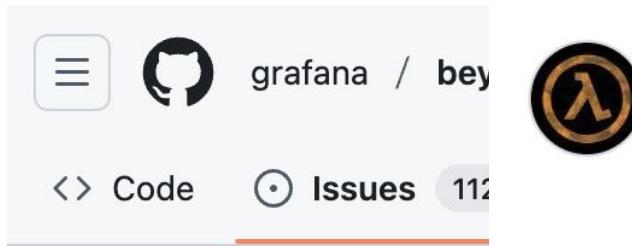


The moment we realised...

- We deployed to a large cluster



Users also reported it



Limit impact of Kubernetes metadata decorator

Closed

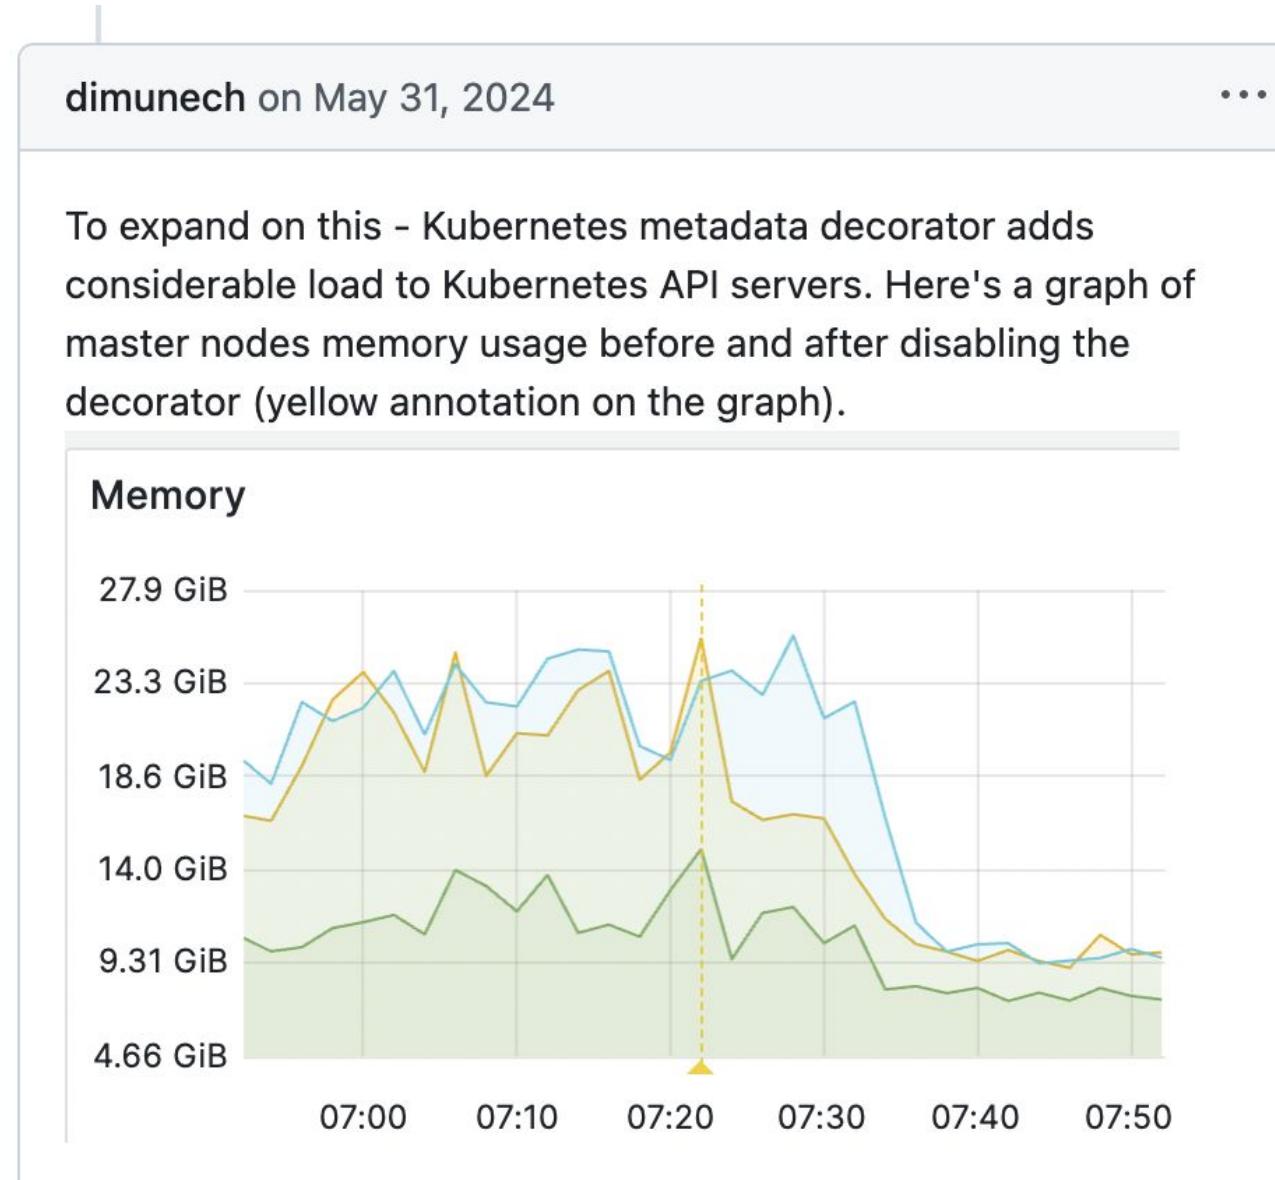


dashpole opened

What I would like

I mentioned this

As a general best practice, when creating replicaset, all services, etc... to watch pods as opposed to apiserver than a deployment.



DDoS'ing Big Clusters

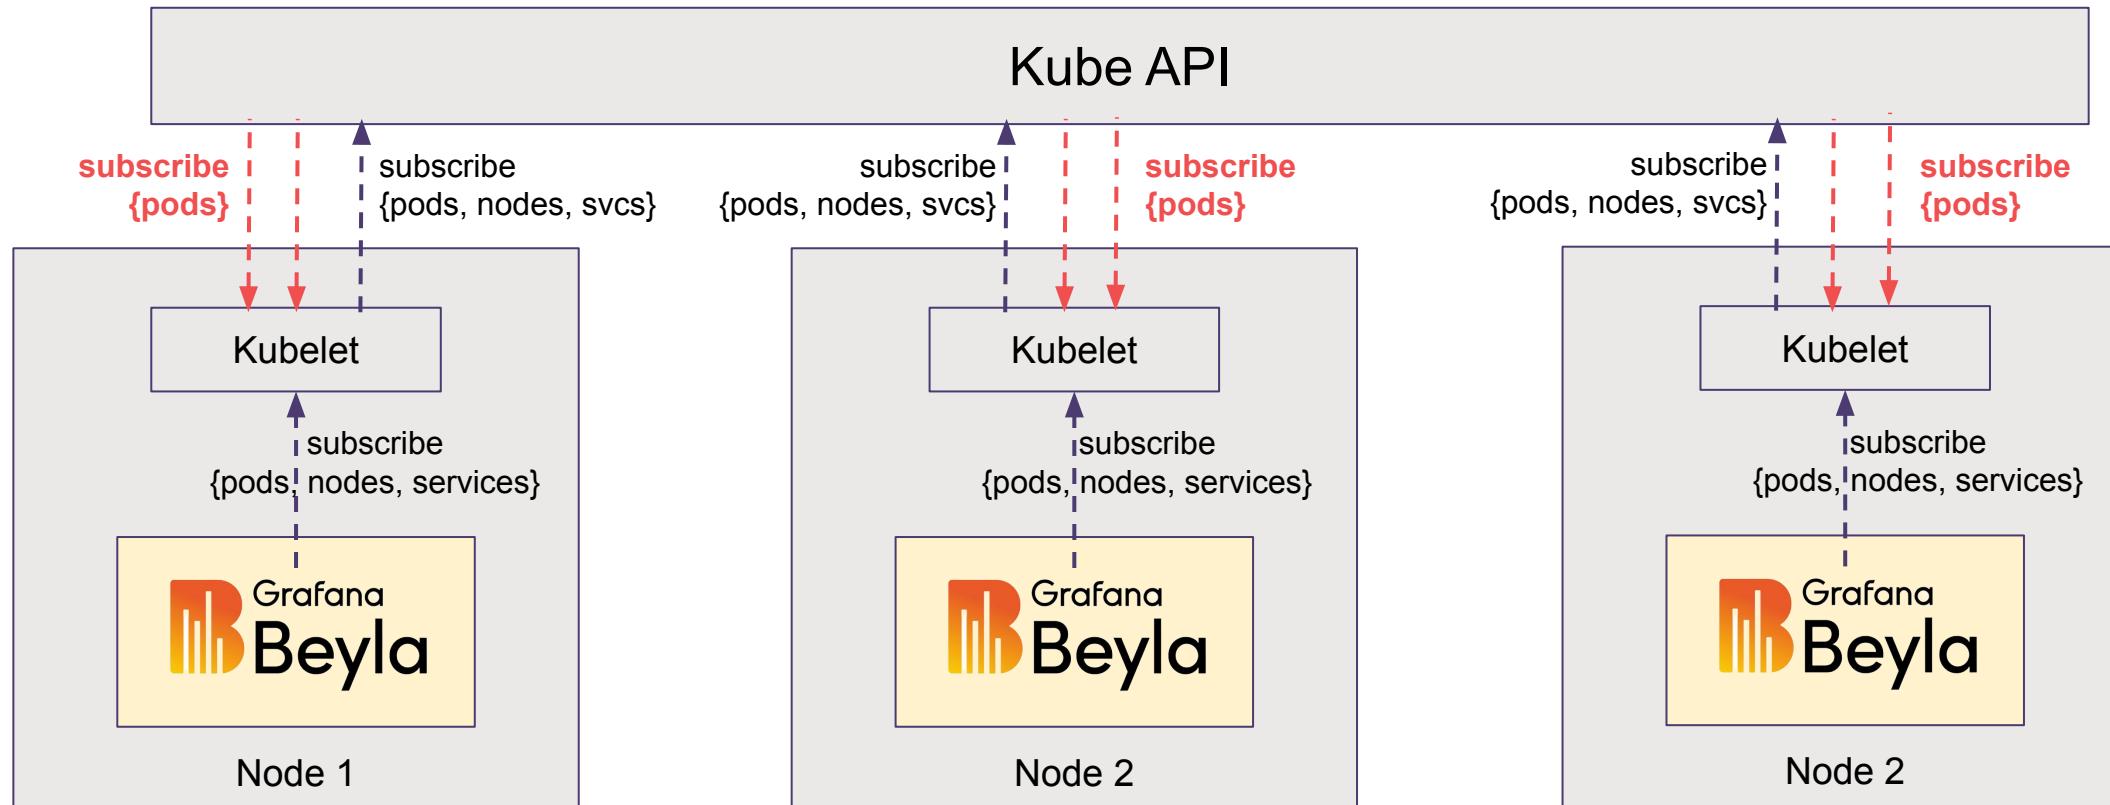


KubeCon
Europe 2025



CloudNativeCon
Europe 2025

Kube API handling $\sim O(N^2)$ subscriptions



How Can We Fix This?

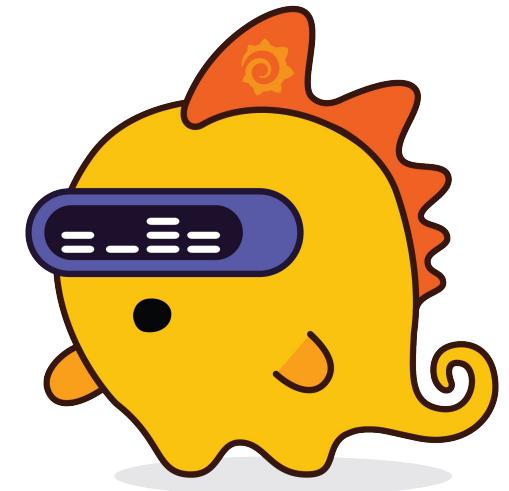
- Replace subscription model by individual requests
- Kubelet API
- Clustered
- Centralized Cache

Replacing subscription model by “get”

- Won't work
 - Need to query information by IP address
- Stampede of requests during deployments

Kubelet API

- The kubelet has an undocumented API
- It runs on every node and maintains its own state
- No access to global objects (e.g. services or Pods from other nodes)

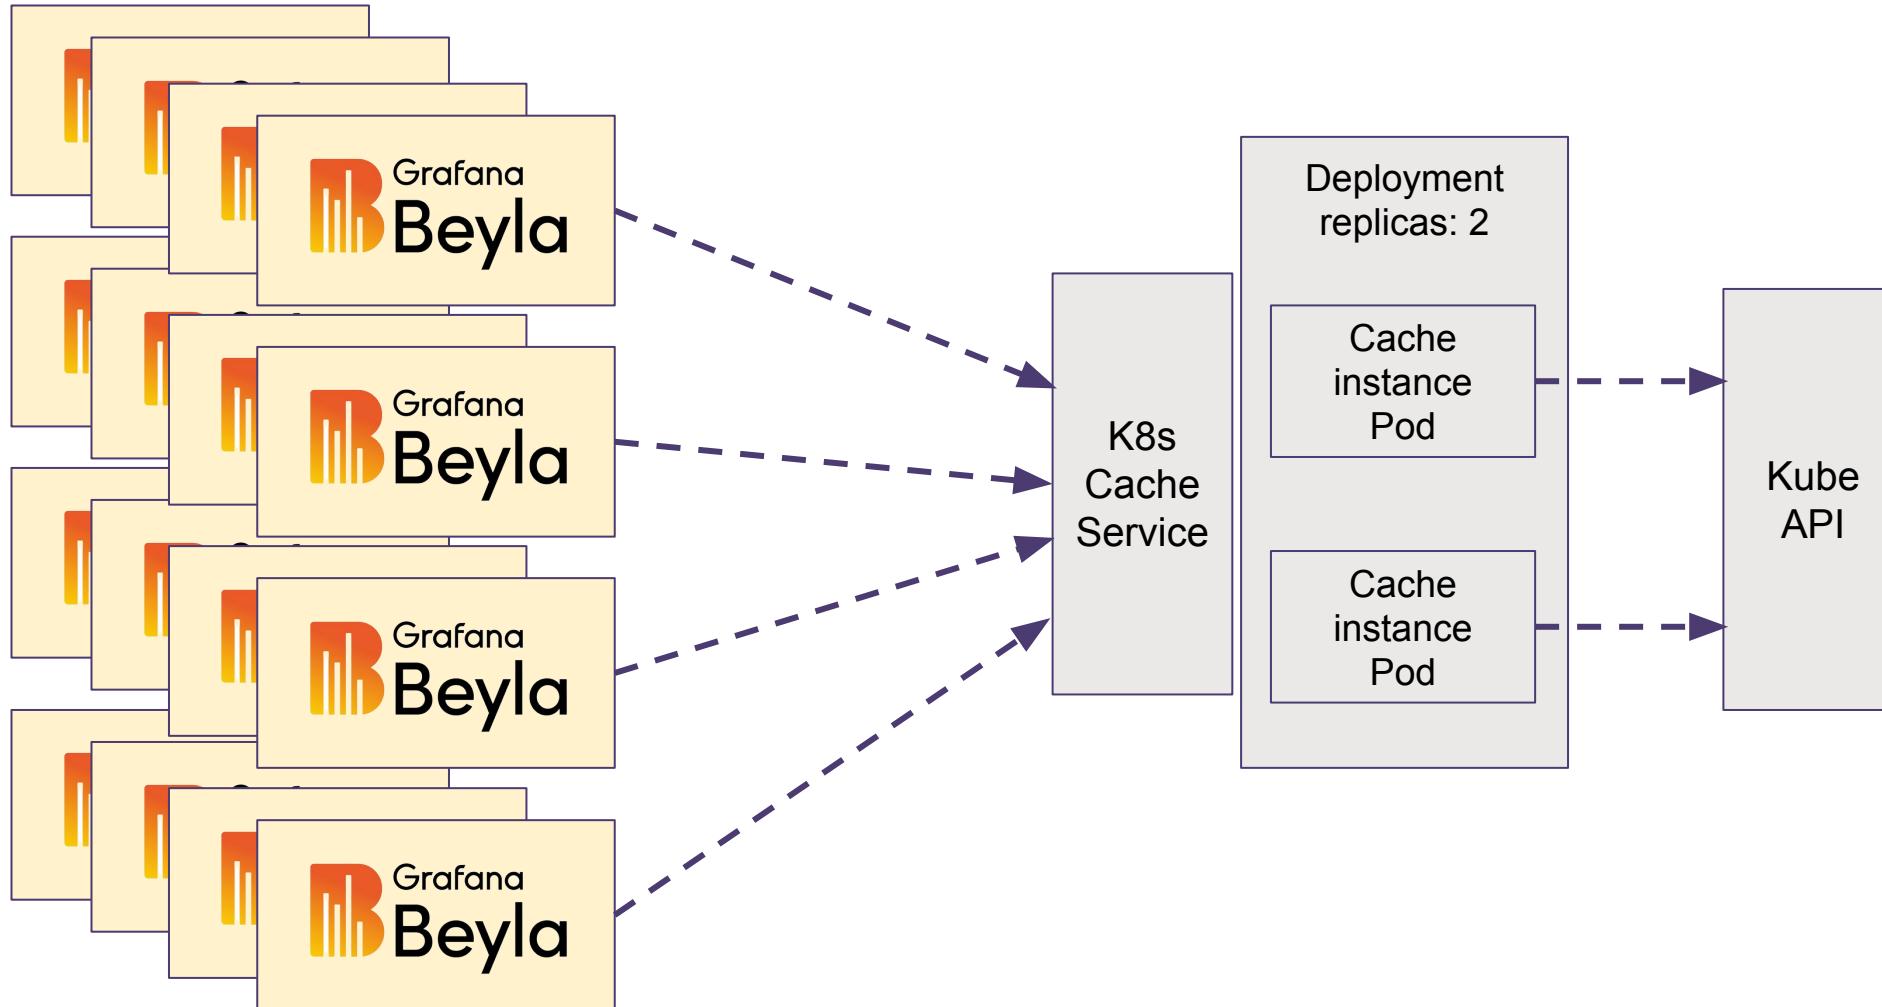


Clustered Cache

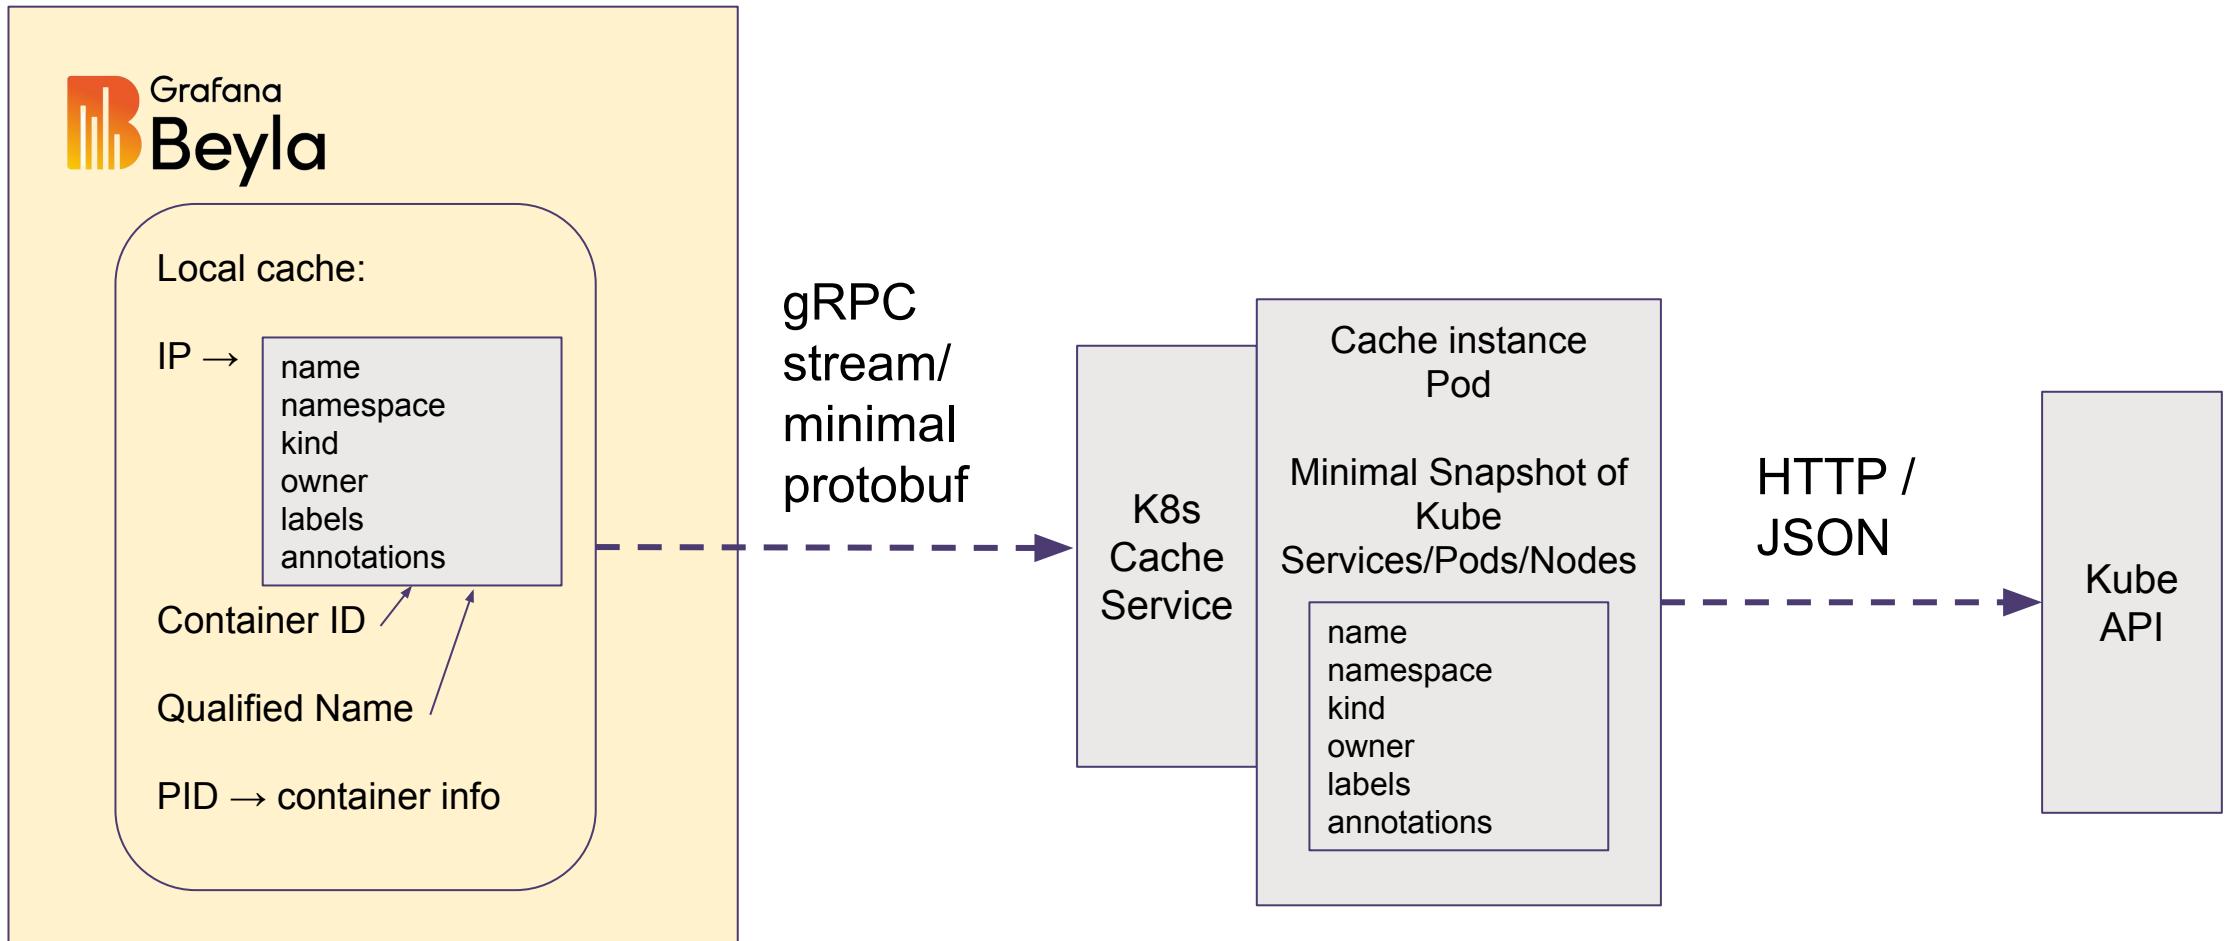
- Using a gossip protocol to share metadata
- Triggered when nodes learn about a new pod <-> IP mappings
- Adds network traffic overhead, complexity

Centralized cache

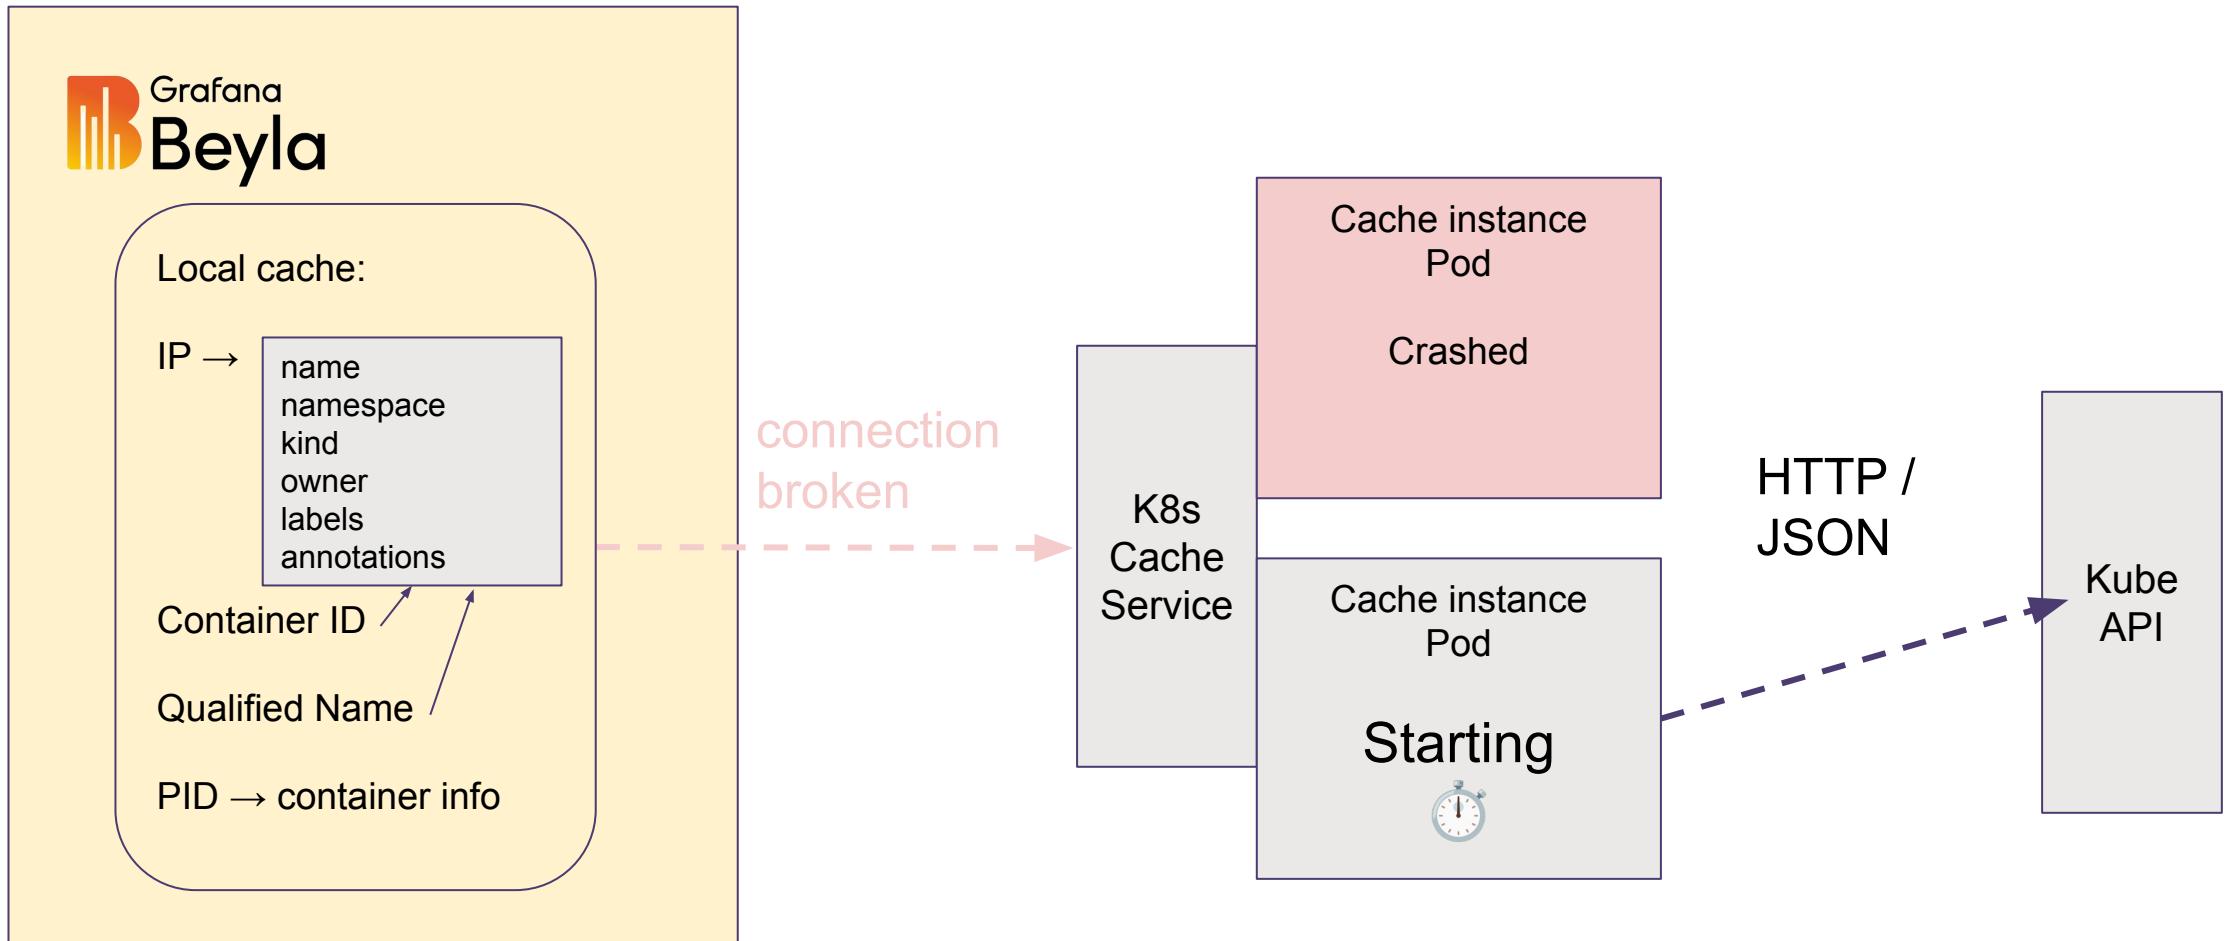
- Chosen approach, for simplicity and flexibility



Centralized cache

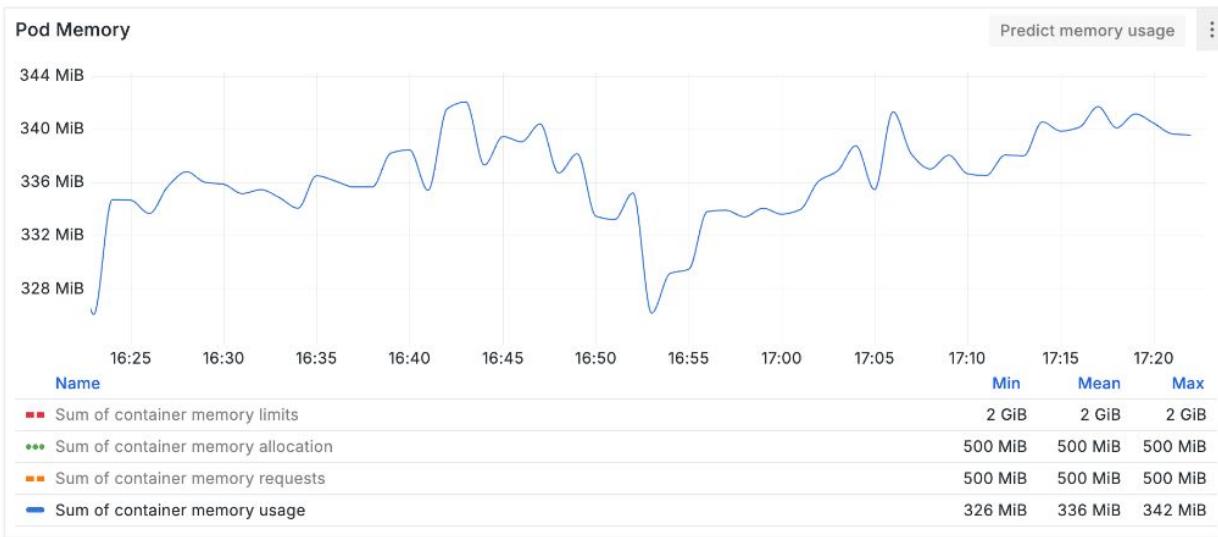
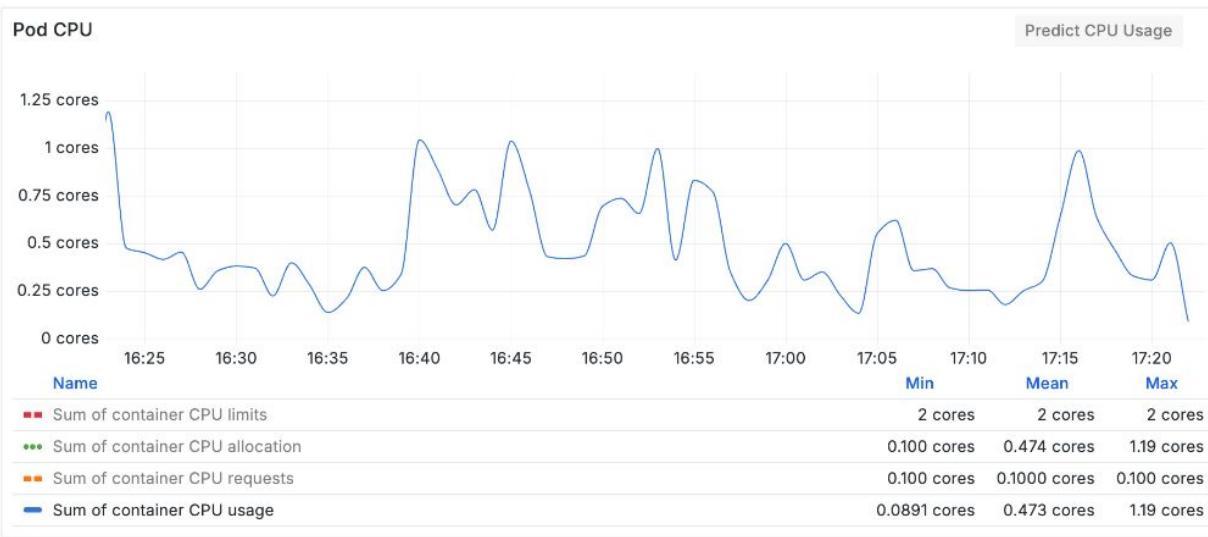


Centralized cache is storageless



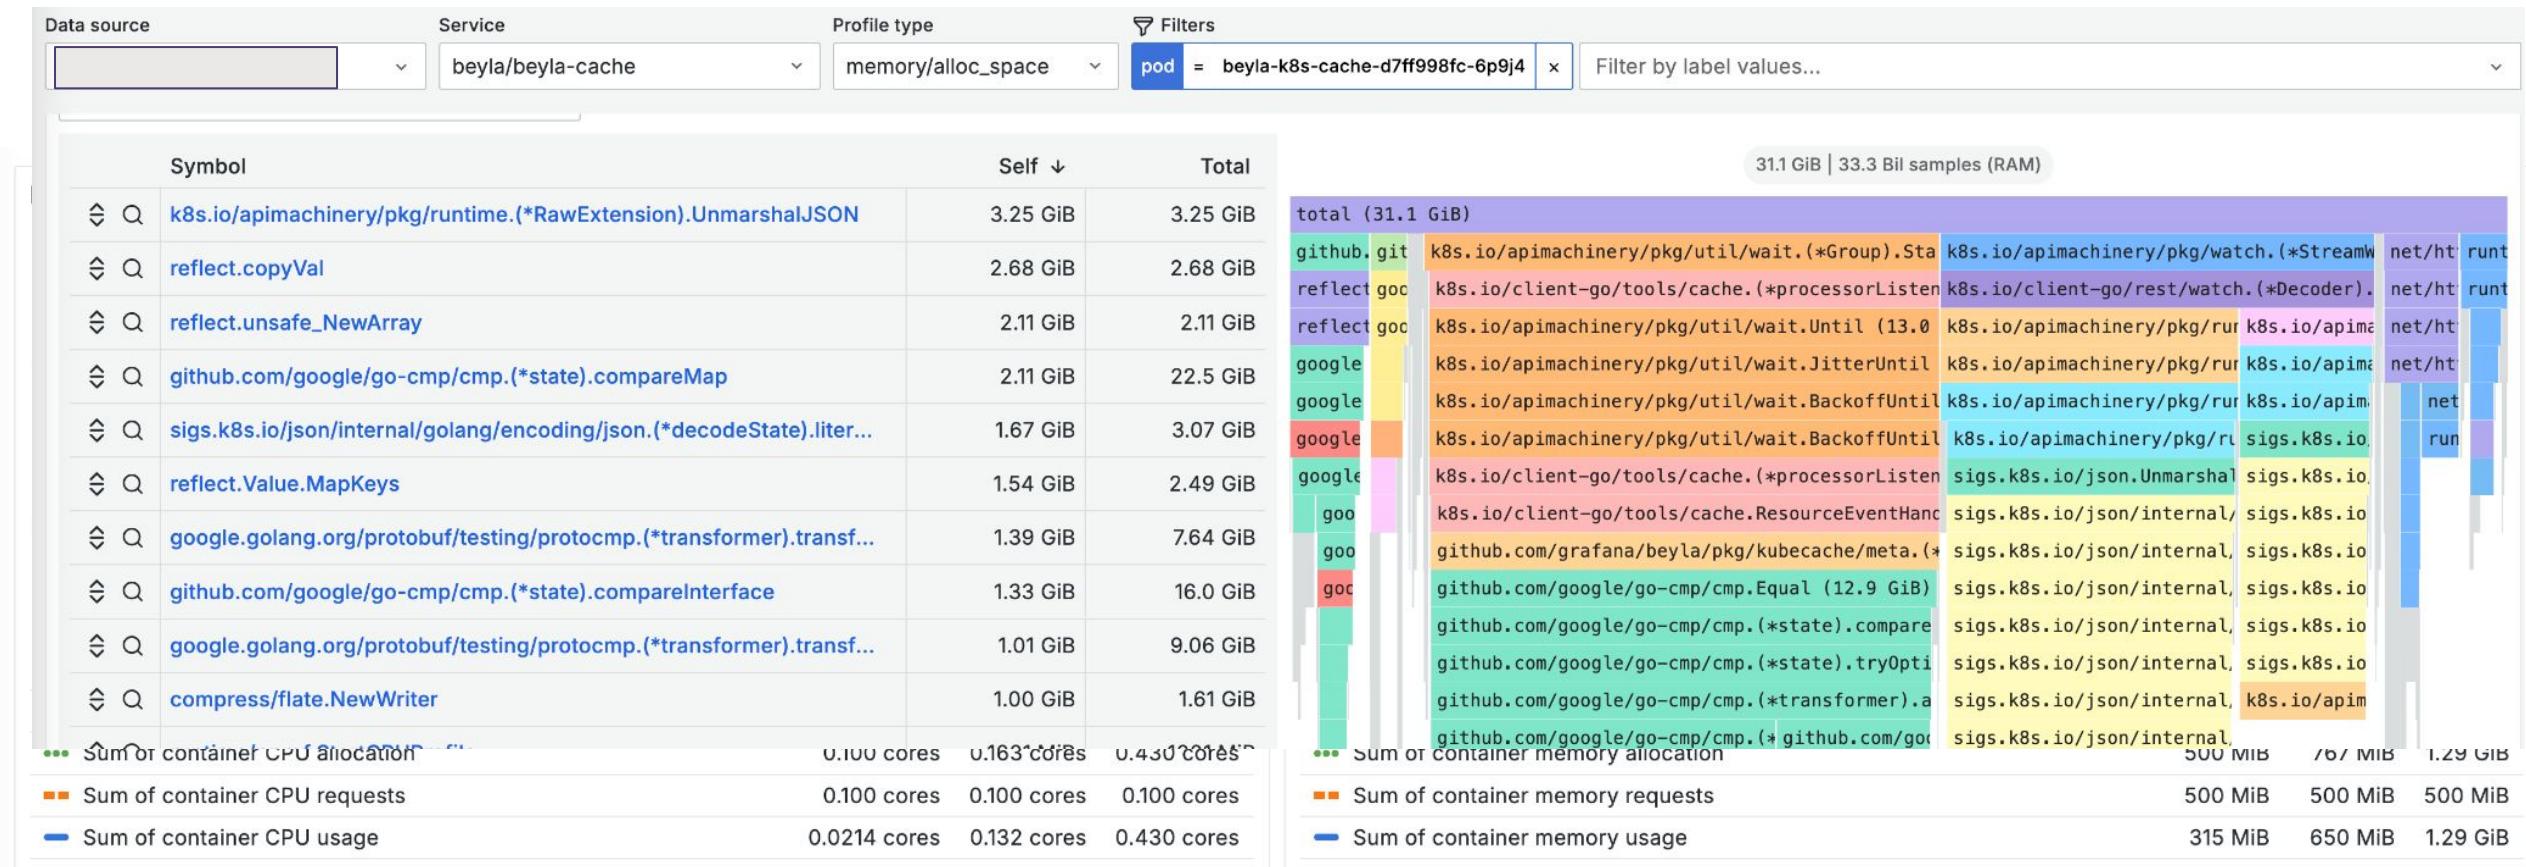
Issues with a centralized cache

- Another component to manage
 - Avoid using external dependencies (DBs, MQs...)
- Resource utilization
 - Especially during startup
 - Use binary encodings
 - Remove unneeded fields



Cache resource utilization

During cache deployment, instances are memory-hungry



Summing Up

- The Kubernetes API can handle a lot, but it has limits
- Our solution was to use a centralized cache
- DaemonSet performance matters a lot

Thank you for your attention!

terra tauri
Staff Software Engineer - Grafana Labs

Mario Macías
Staff Software Engineer - Grafana Labs

